

NTSB Accident Report Analysis

- Author: I Hull
- Date: October 2024
- Title: NTSB Accident Report Analysis USA 2000 to 2020

Source Files:

[NTSB Air Accidents Data](#)

[State Population](#)

[State Name to Abbreviation Lookup](#)

Background

This notebook uses data from the NTSB Air Accidents website for a 24 year period from 01/01/2001 to 01/01/2025.

The NTSB aviation accident database contains civil aviation accidents and selected incidents that occurred from 1962 to present within the United States, its territories and possessions, and in international waters. Foreign investigations in which the NTSB participated as an accredited representative will also be listed.

The data was downloaded as a JSON file and includes 39 fields detailing information relating to the accident or incident such as the date, the aircraft, the type of incident and whether there were any fatalities. After performing some data exploration, this project focuses on the actual report text rather than any other data relating to the incident.

Note: The data downloaded as a CSV seemed to contain different fields such as the make and model of aircraft but for the purposes of this report we will focus on the json file contents that contain the accident reports. Further analysis could be run to include this data and extract further interesting insights.

Research Aim and Motivation

Air incident reports contain a lot of information in summary textual form, including details relating to the conditions at the airport such as weather or runway, pilot actions, the particular conditions relating to the plane etc. This notebook investigates whether distinct topics or themes can be extracted from the text, to provide useful information on the main drivers of air accidents.

Method Notes

The data is cleaned to isolate fields of interest, drop rows containing null values and to extract information for accidents rather than incidents. Air accident: An event that results in serious injury, death, or destruction. Air incident: An event that compromises safety but doesn't result in serious injury, death, or destruction. Detailed method notes are contained within each section of the

notebook, along with any interesting insights extracted from the data from the exploratory data analysis.

Topic Modelling

The data is analysed using Python libraries detailed in the import cell below. Traditional methods of extracting topics from text included Latent Dirichlet Allocation (LDA) and Non-Negative Matrix Factorisation (NMF). These are statistical methods that treat text as a bag of words and they can be complex to implement and interpret.

Newer methods include BERTopic, which uses BERT to create embeddings or representations of text. BERT is short for Bi-directional Encoder Representations, a type of large language model. These representations are clustered into easily interpretable topics. Some text cleaning is necessary to implement this model but less than older topic modelling libraries, since it is designed to work with large-language models which work with context.

1.0 Import libraries

In [125...

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

from wordcloud import WordCloud, STOPWORDS, ImageColorGenerator

import warnings
warnings.filterwarnings("ignore")
```

In [126...

```
# Read in JSON file
df = pd.read_json(r'C:\Users\imoge\AllMLProjects\Data\AirAccidentData.json')
```

In [127...

```
# Shape of dataframe
df.shape
```

Out[127...

(42342, 39)

In [128...

```
# Top two rows
df.head(2)
```

Out[128...

	Oid	MKey	Closed	CompletionStatus	HasSafetyRec	HighestInjury	IsStudy
0	688c61a88126146d93a8e8a7	199500	False	In work	False	None	False
1	688c61a88126146d93a8e8a6	199498	True	Completed	False	None	False

2 rows x 39 columns

In [129...

```
# Columns
df.columns
```

Out[129...

```
Index(['Oid', 'MKey', 'Closed', 'CompletionStatus', 'HasSafetyRec',
      'HighestInjury', 'IsStudy', 'Mode', 'NtsbNumber',
      'OriginalPublishedDate', 'MostRecentReportType', 'ProbableCause',
      'City', 'Country', 'EventDate', 'State', 'Agency', 'BoardLaunch',
      'BoardMeetingDate', 'DocketDate', 'EventType', 'Launch', 'ReportDate',
      'ReportNum', 'ReportType', 'Vehicles', 'AirportId', 'AirportName',
      'AnalysisNarrative', 'FactualNarrative', 'PrelimNarrative',
      'FatalInjuryCount', 'MinorInjuryCount', 'SeriousInjuryCount',
      'InvestigationClass', 'AccidentSiteCondition', 'Latitude', 'Longitude',
      'DocketOriginalPublishDate'],
      dtype='object')
```

In [130...

```
df['EventType'].value_counts(normalize = True)
```

Out[130...

```
ACC    0.937863
INC     0.060602
OCC     0.001535
Name: EventType, dtype: float64
```

In [131...

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 42342 entries, 0 to 42341
Data columns (total 39 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Oid                                    42342 non-null  object
1   MKey                                  42342 non-null  int64
2   Closed                               42342 non-null  bool
3   CompletionStatus                     42342 non-null  object
4   HasSafetyRec                         42342 non-null  bool
5   HighestInjury                       41630 non-null  object
6   IsStudy                              42342 non-null  bool
7   Mode                                 42342 non-null  object
8   NtsbNumber                           42342 non-null  object
9   OriginalPublishedDate                 36687 non-null  object
10  MostRecentReportType                 40180 non-null  object
11  ProbableCause                        35712 non-null  object
12  City                                  42342 non-null  object
13  Country                              42314 non-null  object
14  EventDate                            42342 non-null  object
15  State                                36706 non-null  object
16  Agency                               41122 non-null  object
17  BoardLaunch                          42342 non-null  bool
18  BoardMeetingDate                     7 non-null     object
19  DocketDate                           20 non-null    object
20  EventType                            42342 non-null  object
21  Launch                               20313 non-null  object
22  ReportDate                           33244 non-null  object
23  ReportNum                            93 non-null    object
24  ReportType                           42342 non-null  object
25  Vehicles                             42342 non-null  object
26  AirportId                            26155 non-null  object
```

```

27  AirportName                26155 non-null object
28  AnalysisNarrative          38149 non-null object
29  FactualNarrative           35099 non-null object
30  PrelimNarrative            521 non-null object
31  FatalInjuryCount           42342 non-null int64
32  MinorInjuryCount           42342 non-null int64
33  SeriousInjuryCount         42342 non-null int64
34  InvestigationClass          0 non-null float64
35  AccidentSiteCondition       37297 non-null object
36  Latitude                   38284 non-null float64
37  Longitude                   38284 non-null float64
38  DocketOriginalPublishDate  23642 non-null object
dtypes: bool(4), float64(3), int64(4), object(28)
memory usage: 11.5+ MB

```

Strange that several fields are shown as floats when they are strings

Data Cleaning

In [132...

```

# Check the date column
df['EventDate'].head(2)

```

Out[132...

```

0    2025-01-01T02:20:00Z
1    2024-12-31T14:30:00Z
Name: EventDate, dtype: object

```

In [133...

```

# Check the tail of the date
df['EventDate'].tail(2)

```

Out[133...

```

42340    2001-01-01T16:29:00Z
42341    2001-01-01T13:45:00Z
Name: EventDate, dtype: object

```

In [138...

```

# Set the event date to a datetime object
df['EventDate'] = pd.to_datetime(df['EventDate'])

```

In [139...

```

# Create a year and month column from the date
df['Year'] = df['EventDate'].dt.year
df['Month'] = df['EventDate'].dt.month

# Set to string
df['Year'] = df['Year'].astype('object')
df['Month'] = df['Month'].astype('object')

```

In [140...

```

# Check for nulls
df.isnull().sum()

```

Out[140...

```

Oid                0
MKey               0
Closed             0
CompletionStatus   0
HasSafetyRec       0
HighestInjury      712
IsStudy            0
Mode              0
NtsbNumber         0

```

OriginalPublishedDate	5655
MostRecentReportType	2162
ProbableCause	6630
City	0
Country	28
EventDate	0
State	5636
Agency	1220
BoardLaunch	0
BoardMeetingDate	42335
DocketDate	42322
EventType	0
Launch	22029
ReportDate	9098
ReportNum	42249
ReportType	0
Vehicles	0
AirportId	16187
AirportName	16187
AnalysisNarrative	4193
FactualNarrative	7243
PrelimNarrative	41821
FatalInjuryCount	0
MinorInjuryCount	0
SeriousInjuryCount	0
InvestigationClass	42342
AccidentSiteCondition	5045
Latitude	4058
Longitude	4058
DocketOriginalPublishDate	18700
Year	0
Month	0
dtype: int64	

```
In [141... # How many were recorded as accidents and how many as incidents?
df['EventType'].value_counts()
```

```
Out[141... ACC      39711
INC       2566
OCC         65
Name: EventType, dtype: int64
```

```
In [142... # Drop incidents and focus on accidents
df = df[df['EventType']=='ACC']
df.shape
```

```
Out[142... (39711, 41)
```

```
In [143... # Choose columns of interest
df = df[['ProbableCause', 'City', 'EventDate', 'Year', 'Month', 'State',
        'FactualNarrative', 'FatalInjuryCount', 'MinorInjuryCount',
        'SeriousInjuryCount', 'HighestInjury', 'AccidentSiteCondition',
        'Latitude', 'Longitude']]
```

```
In [144... df.head()
```


Out[144...

	ProbableCause	City	EventDate	Year	Month	State	FactualNarrative	FatalInjuryCount	M
0	None	Naples	2025-01-01 02:20:00+00:00	2025	1	FL	None	0	
1	The student pilot's failure to maintain direct...	Anchorage	2024-12-31 14:30:00+00:00	2024	12	AK	None	0	
2	The pilot's failure to maintain airplane contr...	Los Alamos	2024-12-31 14:15:00+00:00	2024	12	NM	None	0	
3	The pilot's failure to maintain terrain cleara...	Galveston	2024-12-31 13:20:00+00:00	2024	12	TX	None	0	
4	None	Peebles	2024-12-30 16:15:00+00:00	2024	12	OH	None	1	



Data Exploration

In [145...

```
df.columns
```

Out[145...

```
Index(['ProbableCause', 'City', 'EventDate', 'Year', 'Month', 'State',  
      'FactualNarrative', 'FatalInjuryCount', 'MinorInjuryCount',  
      'SeriousInjuryCount', 'HighestInjury', 'AccidentSiteCondition',  
      'Latitude', 'Longitude'],  
      dtype='object')
```

In [146...

```
df.info()  
  
<class 'pandas.core.frame.DataFrame'>  
Int64Index: 39711 entries, 0 to 42341  
Data columns (total 14 columns):  
#   Column                                Non-Null Count  Dtype  
---  -  
0   ProbableCause                        34858 non-null  object  
1   City                                39711 non-null  object  
2   EventDate                           39711 non-null  datetime64[ns, UTC]  
3   Year                                39711 non-null  object  
4   Month                               39711 non-null  object  
5   State                               35646 non-null  object  
6   FactualNarrative                    33432 non-null  object  
7   FatalInjuryCount                    39711 non-null  int64  
8   MinorInjuryCount                    39711 non-null  int64  
9   SeriousInjuryCount                  39711 non-null  int64  
10  HighestInjury                        39386 non-null  object  
11  AccidentSiteCondition                36389 non-null  object  
12  Latitude                             37017 non-null  float64  
13  Longitude                            37018 non-null  float64  
dtypes: datetime64[ns, UTC](1), float64(2), int64(3), object(8)  
memory usage: 4.5+ MB
```

We have a few null values for latitude and longitude but we won't worry about those at this point

Accidents by Year

In [149...

```
year_accidents = pd.DataFrame(df['Year'].value_counts().sort_index())
year_accidents = year_accidents.iloc[0:-1]
```

In [162...


```
plt.figure(figsize = (14,6))
ax = sns.lineplot(data = year_accidents)

labels = year_accidents.index.astype('str')

#calculate equation for trendline
z = np.polyfit(year_accidents.index,year_accidents.Year.values, 1)
p = np.poly1d(z)

#add trendline to plot
plt.plot(year_accidents.index, p(year_accidents.index), color = 'k',linestyle='dashed')

# Title and Legend
plt.title("Air Accidents USA 01/01/2001 to 31/12/2024", fontsize = 15)
ax.spines[['right', 'top']].set_visible(False)
ax.legend("")
ax.set_xticks(year_accidents.index)
plt.text(0.67,0.086,'COVID Lockdown', transform=ax.transAxes, color = 'firebrick')
ax.set_xticklabels(labels);
```



Trending down over the period and sharp fall in 2020 during Covid lockdown

Accidents by Month

In [267...

```
month_accidents = df.groupby('Month',as_index = False)['Year'].count()
month_accidents.columns = ['Month','Number']
month_accidents
```

Out[267...

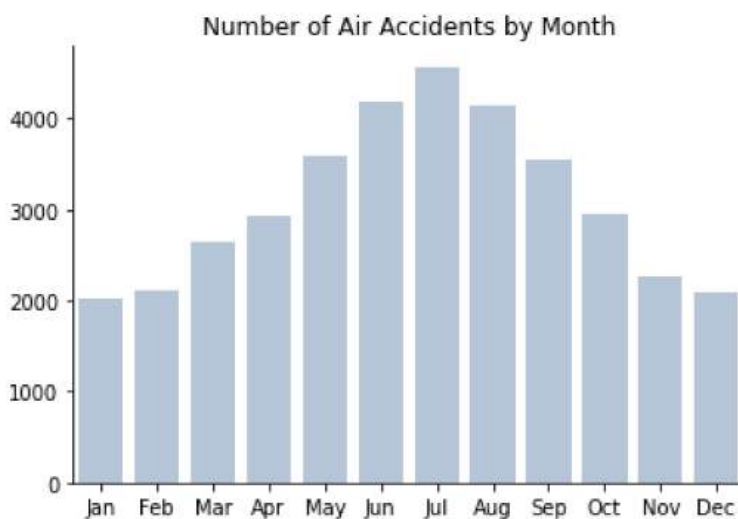
	Month	Number
0	1	2021
1	2	2104
2	3	2646
3	4	2928
4	5	3582
5	6	4192
6	7	4565
7	8	4131
8	9	3535
9	10	2959
10	11	2273
11	12	2080

In [268...

```
# Replace month numbers with month names (abbreviated or full)
import calendar
month_accidents['Month'] = month_accidents['Month'].apply(lambda x: calendar.month_abbr[x])

month_order = list(calendar.month_abbr)[1:] # ['Jan', 'Feb', ..., 'Dec']
month_accidents['Month'] = pd.Categorical(month_accidents['Month'], categories=month_order)

ax = sns.barplot(data = month_accidents, x = 'Month', y = 'Number', color='lightsteelblue')
ax.spines[['right', 'top']].set_visible(False)
plt.ylabel("")
plt.xlabel("")
plt.title("Number of Air Accidents by Month");
```



Accidents by Injury

In [163...

```
injuries = pd.DataFrame(df['HighestInjury'].value_counts(normalize = True)*100)
injuries
```


Out[163...

HighestInjury	
None	51.614787
Fatal	22.335348
Minor	14.355355
Serious	11.694511


In [164...

```
#define data
data = injuries.HighestInjury
labels = injuries.index

#define Seaborn color palette to use
colors = sns.color_palette('mako')[0:5]

#create pie chart
plt.pie(data,
        labels = labels,
        colors = colors,
        autopct='%.0f%%',
        textprops={'color':"w", 'fontsize': 14})
plt.title("Injuries By Type", fontsize = 15)
plt.legend(bbox_to_anchor=(1, 0.7))
plt.show()
```

Injuries By Type



In [165...

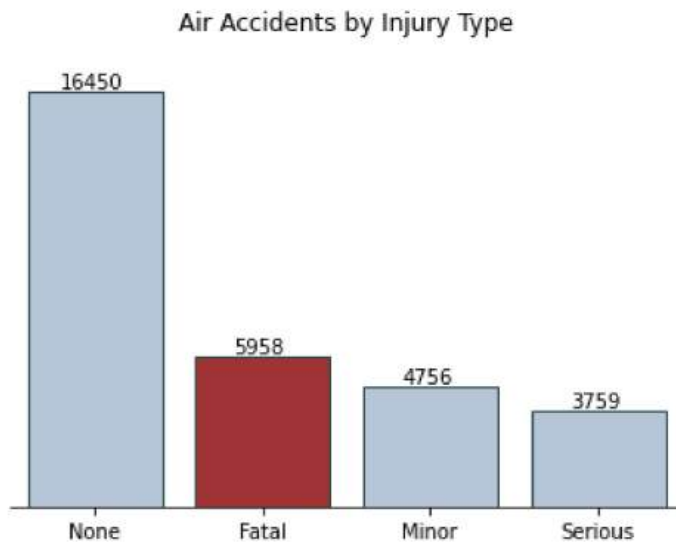
```
injuries = pd.DataFrame(df3['HighestInjury'].value_counts())
injuries
```

Out[165...

HighestInjury	
None	16450
Fatal	5958
Minor	4756
Serious	3759

In [167...

```
colors = ['lightsteelblue', 'firebrick', 'lightsteelblue', 'lightsteelblue'] # 'red' f
ax = sns.barplot(data = injuries, x = injuries.index,
                 y = 'HighestInjury',
                 palette = colors,
                 ec = 'darkslategray')
ax.spines[['right', 'top', 'left']].set_visible(False)
plt.title("Air Accidents by Injury Type", fontsize = 12, pad = 20)
plt.ylabel("Number", labelpad = 10)
plt.yticks([])
plt.ylabel("")
# Adding values on top of the bars
for p in ax.patches:
    ax.annotate(format(p.get_height(), '.0f'),
                (p.get_x() + p.get_width() / 2.1, p.get_height()),
                ha='center', va='bottom');
```



- Most accidents have no injuries
- When there are injuries the majority of fatal

In [168...

```
# Fatal injuries analysis
fatal = pd.DataFrame(df[df['FatalInjuryCount']>0]['FatalInjuryCount'].value_counts()).h
fatal.columns = ['Fatal Injury Count']
fatal['%'] = round(fatal['Fatal Injury Count']/fatal['Fatal Injury Count'].sum()*100,1)
fatal
```

Out[168...

	Fatal Injury Count	%
1	4453	51.4
2	2511	29.0
3	765	8.8
4	486	5.6
5	198	2.3
6	99	1.1
7	61	0.7

	Fatal Injury Count	%
9	34	0.4
8	32	0.4
10	25	0.3

In [175...

```
pd.DataFrame(df[df['FatalInjuryCount']>0]['FatalInjuryCount'].value_counts()).sort_index
```

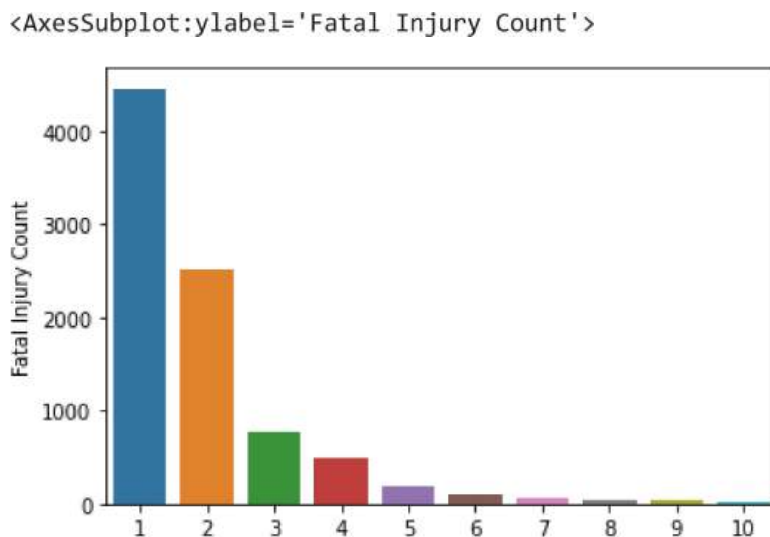
Out[175...

	FatalInjuryCount
298	1
265	1
228	1
224	1
206	1

In [170...

```
#create chart
sns.barplot(data = fatal, x = fatal.index, y = fatal['Fatal Injury Count'])
```

Out[170...



- Most fatal accidents involve 1 or 2 people and very few involve many
- In the 20 year period only one accident involved large numbers of people

In [171...

```
# Fatal accidents with more than 20 people involved
pd.DataFrame(df[df['FatalInjuryCount']>0]['FatalInjuryCount'].value_counts()).sort_index
```

Out[171...

	FatalInjuryCount
298	1
265	1
228	1

FatalInjuryCount	
224	1
206	1
199	1
191	1
176	1
173	1
162	1

In [177...

```
# Find details on this flight
df[df['FatalInjuryCount']==265]
```

Out[177...

	ProbableCause	City	EventDate	Year	Month	State	FactualNarrative	FatalInjuryCount	N
40499	the in-flight separation of the vertical stabi...	Belle Harbor	2001-11-12 10:16:00+00:00	2001	11	NY	The Board's full report is available at http://...	265	



American Airlines Flight 587 was a regularly scheduled international passenger flight from John F. Kennedy International Airport, New York City to Las Américas International Airport, Santo Domingo. On November 12, 2001, the Airbus A300B4-605R flying the route crashed into the neighborhood of Belle Harbor on the Rockaway Peninsula of Queens, New York City, shortly after takeoff, killing all 260 people aboard (251 passengers and 9 crew members), as well as 5 people on the ground.[1] It is the second-deadliest aviation accident in U.S. history, behind the crash of American Airlines Flight 191 in 1979,[a][1] and the second-deadliest aviation incident involving an Airbus A300, after Iran Air Flight 655.[1][3]

Accidents by Site Condition

In [178...

```
conditions = pd.DataFrame(df['AccidentSiteCondition'].value_counts(normalize = True)*100)
conditions
```

Out[178...

AccidentSiteCondition	
VMC	93.816813
IMC	5.531892
Unknown	0.651296

VMC = Visual Meterological Conditions

IMC = Instrument Meterological Conditions (Poor weather, Flying in clouds, Limited visibility)


In [179...

```
#define data
data = conditions.AccidentSiteCondition[0:2]
labels = conditions.index[0:2]

#define Seaborn color palette to use
colors = ['lightsteelblue', 'slategray'] # Replace with your desired colors

#create pie chart
plt.pie(data,
        labels = labels,
        colors = colors,
        autopct='%0f%%',
        textprops={'color': "w", 'fontsize': 14})
plt.title("Accidents by Flying Conditions", fontsize = 12)
plt.legend(bbox_to_anchor=(1, 0.7))
plt.show()
```

Accidents by Flying Conditions



Accidents by State

In [181...

```
# Accidents by state
state_accidents = pd.DataFrame(df['State'].value_counts()).reset_index()
state_accidents.columns = ['State', 'Accidents']
state_accidents['PerYear'] = round(state_accidents['Accidents']/24,0)
state_accidents.head()
```

Out[181...

	State	Accidents	PerYear
0	CA	3394	141.0
1	TX	2747	114.0
2	FL	2589	108.0
3	AK	2326	97.0
4	AZ	1412	59.0

In [182...

```
# Get a state name and abbreviation lookup table
abbrev = pd.read_excel(r'C:\Users\imoge\AllMLProjects\Data\Abbreviations.xlsx')
abbrev.columns = ['Name', 'Abbrev']
```



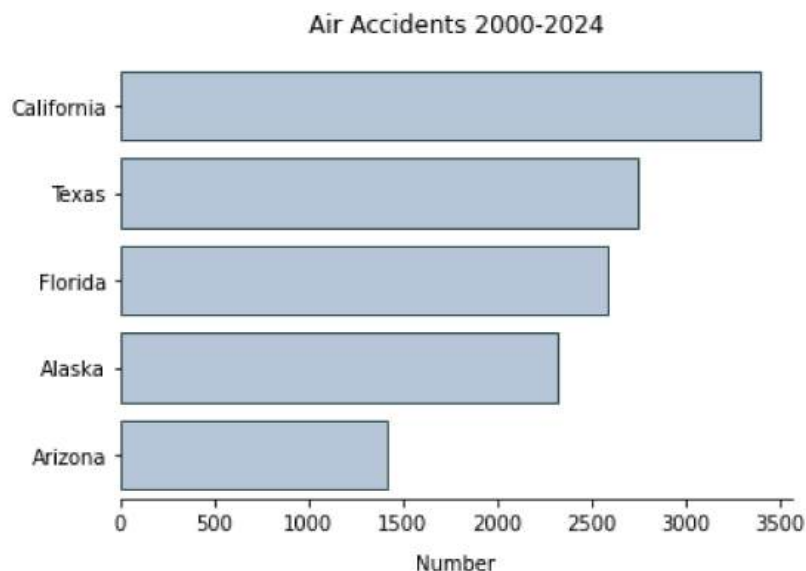
```
# Merge with the accident table to give names to the states
state_accidents = state_accidents.merge(abrev, how = 'left', left_on = 'State', right_o
state_accidents.drop(columns = 'Abrev',axis = 1, inplace = True)
state_accidents.head()
```

Out[182...

	State	Accidents	PerYear	Name
0	CA	3394	141.0	California
1	TX	2747	114.0	Texas
2	FL	2589	108.0	Florida
3	AK	2326	97.0	Alaska
4	AZ	1412	59.0	Arizona

In [184...

```
# Plot the air accidents by state in total for 2000 to 2020
ax = sns.barplot(data = state_accidents[0:5], y = 'Name',
                 x = 'Accidents',
                 color = 'lightsteelblue',
                 ec = 'darkslategray')
ax.spines[['right', 'top', 'left']].set_visible(False)
ax.set_ylabel("")
ax.set_xlabel("Number", labelpad = 10)
ax.set_title("Air Accidents 2000-2024", fontsize = 12, pad = 15);
```



Most accidents are reported for California, followed by Texas, Florida. Most populous states so expect more flights

In [185...

```
# Get average state population from 2020 and 2023
pop = pd.read_excel(r'C:\Users\imoge\AllMLProjects\Data\USPopulation.xlsx')
pop['AveragePop'] = (pop['Pop. 2000'] + pop['Pop. 2023'])/2
pop = pop[['Name', 'AveragePop']]
pop.head()
```

Out[185...

	Name	AveragePop
0	United States	308169747.5
1	Alabama	4777837.5
2	Alaska	680169.5
3	Arizona	6280795.5
4	Arkansas	2870512.5

In [186...

```
# Merge the abbreviation table with the state accidents table
combined = state_accidents.merge(pop, how = 'left', left_on = 'Name', right_on = 'Name')
combined.head()
```

Out[186...

	State	Accidents	PerYear	Name	AveragePop
0	CA	3394	141.0	California	36418423.0
1	TX	2747	114.0	Texas	25677164.5
2	FL	2589	108.0	Florida	19296648.5
3	AK	2326	97.0	Alaska	680169.5
4	AZ	1412	59.0	Arizona	6280795.5

In [187...

```
# Calculate the average annual accidents per million population
combined['AccperYearperMill'] = round(combined['PerYear']/combined['AveragePop']*100000)
combined.sort_values(by = 'AccperYearperMill', ascending = False).head()
```

Out[187...

	State	Accidents	PerYear	Name	AveragePop	AccperYearperMill
3	AK	2326	97.0	Alaska	680169.5	142.6
38	WY	314	13.0	Wyoming	538921.5	24.1
9	ID	805	34.0	Idaho	1629341.5	20.9
29	MT	476	20.0	Montana	1017506.0	19.7
22	NM	587	24.0	New Mexico	1966694.0	12.2

In [188...

```
combined[combined['Name']=='California']
```

Out[188...

	State	Accidents	PerYear	Name	AveragePop	AccperYearperMill
0	CA	3394	141.0	California	36418423.0	3.9

In [189...

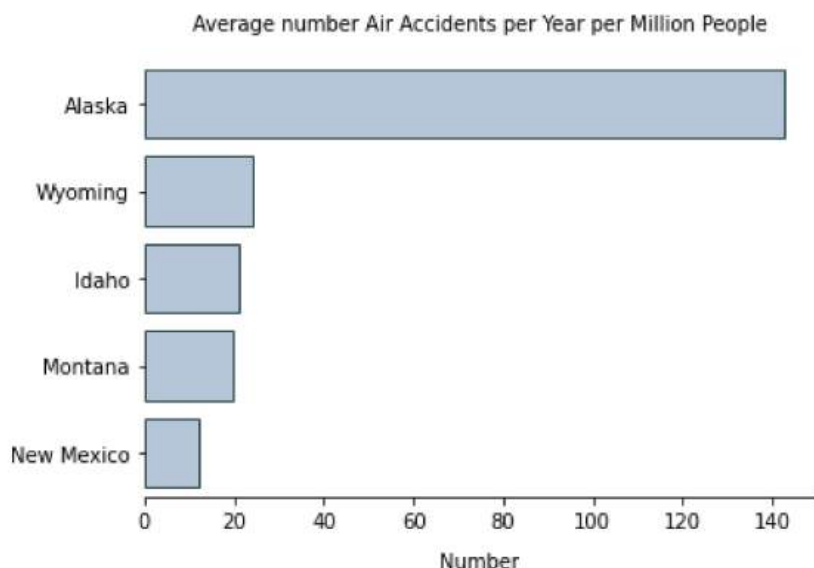
```
# Top five
top_five = combined.sort_values(by = 'AccperYearperMill', ascending = False).head(5)
top_five
```

```
Out[189...
State Accidents PerYear Name AveragePop AccperYearperMill
3 AK 2326 97.0 Alaska 680169.5 142.6
38 WY 314 13.0 Wyoming 538921.5 24.1
9 ID 805 34.0 Idaho 1629341.5 20.9
29 MT 476 20.0 Montana 1017506.0 19.7
22 NM 587 24.0 New Mexico 1966694.0 12.2
```

```
In [190...
# Bottom five
bottom_five = combined.sort_values(by = 'AccperYearperMill',ascending = True).head()
bottom_five
```

```
Out[190...
State Accidents PerYear Name AveragePop AccperYearperMill
55 DC 4 0.0 District of Columbia 625529.0 0.0
48 PR 113 5.0 Puerto Rico 3507148.0 1.4
12 NY 758 32.0 New York 19274121.0 1.7
50 RI 57 2.0 Rhode Island 1072110.5 1.9
36 MA 320 13.0 Massachusetts 6675381.5 1.9
```

```
In [191...
ax = sns.barplot(data = top_five, y = 'Name',
                 x = 'AccperYearperMill',
                 color = 'lightsteelblue',
                 ec = 'darkslategray')
ax.spines[['right', 'top','left']].set_visible(False)
ax.set_ylabel("")
ax.set_xlabel("Number", labelpad = 10)
ax.set_title("Average number Air Accidents per Year per Million People",fontsize = 10,
```




```
In [98]:
# Split off data of interest
#acc_mill = top_ten[['Name','AccperYearperMill']]
```

```
#acc_mill.columns = [['State', 'AccperYearMill']]
#acc_mill
```

Accidents by City

```
In [192... df['City'].value_counts().head(10).plot(kind = 'barh')
```

Out[192... <AxesSubplot:>



Location of Accidents

```
In [193... # Remove rows where we have no data on Latitude and Longitude
df = df[~df['Latitude'].isnull()]
df = df[~df['Longitude'].isnull()]
```

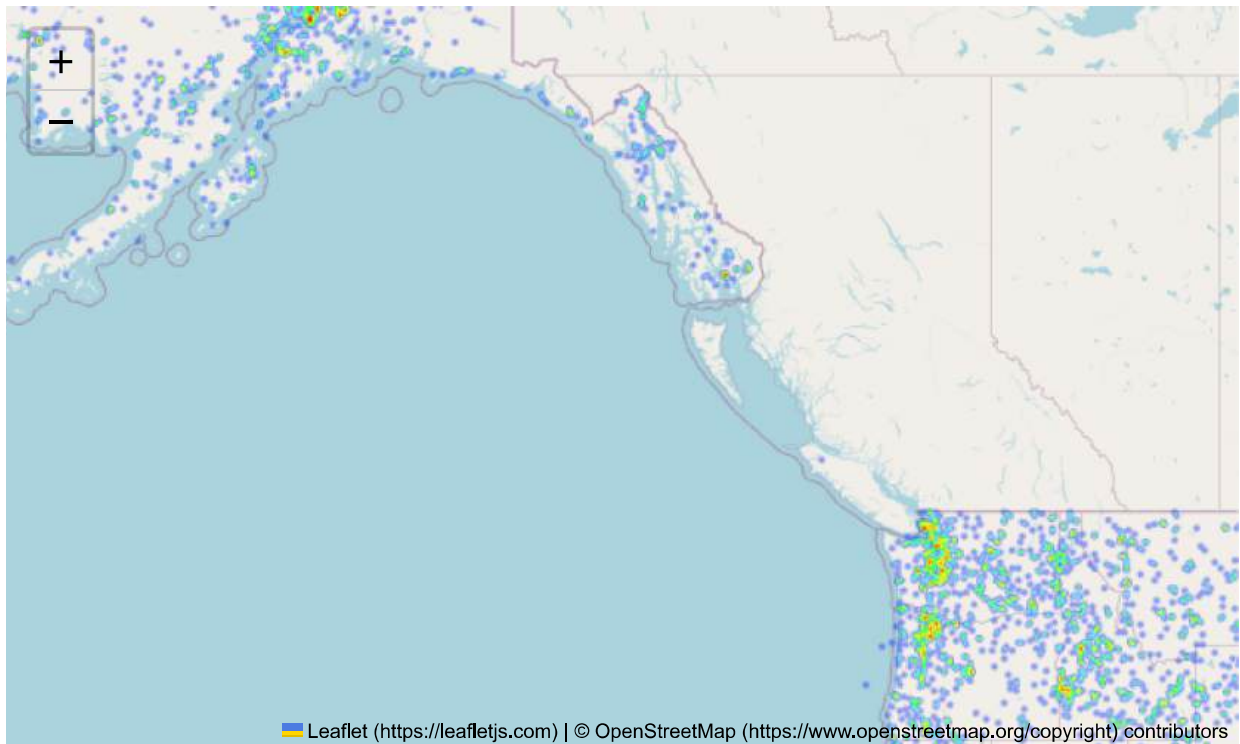
```
In [194... from mpl_toolkits.basemap import Basemap
import folium
from folium import plugins
```

```
In [195... # Convert Latitude and Longitude to Lists
lat = df3['Latitude'].to_list()
long = df3['Longitude'].to_list()

# Create map centred on the US with a zoom
accident_map = folium.Map([38.9, -103], zoom_start=4.49)

heatmap = plugins.HeatMap(list(zip(lat, long)), radius = 2, blur = 1)
accident_map.add_child(heatmap)
```

Out[195...



Text Preprocessing

In [205...

```
# Find the reports that have completed
df2 = df[~df['FactualNarrative'].isnull()]
df2.shape
```

Out[205...

(32060, 14)

In [206...

```
# Look at summary of an example report
df2['ProbableCause'].iloc[3]
```

Out[206...

'The pilot's improper fuel planning, which resulted in fuel exhaustion and a total loss of power to both engines.'

In [209...

```
df2.iloc[3]['FactualNarrative']
```

Out[209...

'On November 27, 2024, about 0005 central standard time, a Cessna 340A airplane, N5757C, sustained substantial damage when it was involved in an accident near Muskogee, Oklahoma. The pilot and one passenger sustained minor injuries, and three passengers were not injured. The airplane was operated as a Title 14 Code of Federal Regulations Part 91 personal flight. The pilot reported that, before departing on the cross-country flight, the airplane was "topped off" with 100 low lead fuel, resulting in a total of 183 gallons onboard for the flight. The pilot performed a preflight inspection of the airplane at night, but he did not have his headlamp with him, only his cellular phone. The pilot reported that, when verifying fuel levels in the dark, he typically feels inside the fuel tank with his finger, but he recalled not doing that with the right main fuel tank before departing on the accident flight. The pilot's fuel calculations showed that the flight would require a total of 150 gallons of fuel. The airplane was equipped with two main fuel tanks, each with a total capacity of 51 gallons, of which 50 gallons was usable; two auxiliary fuel tanks, each with a capacity of 32 gallons (31.5 gallons usable each tank); and one left side locker fuel tank, which had a total capacity of 20.5 gallons (20 gallons usable). A review of ADS-B data showed that the airplane depa

rted from Mission Field Airport (LVM), Livingston, Montana, about 1840 on an instrument flight rules flight plan to the intended destination of Muskogee-Davis Regional Airport (MKO), Muskogee, Oklahoma.

The pilot reported a higher than usual fuel consumption rate during the first half of the flight, which was in instrument meteorological conditions (IMC). While operating in IMC, the pilot was operating the engines at 100° degrees rich of peak versus the normal peak engine gas temperature. Additionally, the pilot was operating the heater for the first half of the flight and then occasionally for the second half of the flight. The pilot reported that, while enroute, he transferred fuel from the auxiliary fuel tanks and the left side locker fuel tank to the main fuel tanks. After the "final transfer" was completed, the pilot calculated that there was adequate fuel remaining, including the required reserve fuel, to make the destination airport. However, he reported that, over Tulsa (about 50 miles from the destination airport), the indicated fuel levels were less than he expected.

Shortly after the pilot was cleared for a visual approach at the destination, the right engine lost total power, followed by the left engine. The pilot attempted to restart both engines to no avail. At the beginning of the loss of engine power sequence, the pilot reported that the left side of the airplane had about 30 lbs (about 4.84 gallons) of fuel remaining and the right side had about 30 to 35 lbs (about 4.84 to 5.65 gallons) of fuel remaining.

During the subsequent forced landing, the right wing impacted a permanent elevated static display on MKO property, a U.S. Air Force T-33A jet trainer, located about 3,350 ft northwest of the approach end of runway 13, and came to rest upright. The pilot and passengers egressed the airplane without further incident. The airplane sustained substantial damage to the fuselage and the right wing.

During an onsite examination, fuel blighting was not observed on the grass. The right main fuel tank and the right auxiliary fuel tank were separated from the right wing during the impact sequence and sustained damage. All other fuel tanks were found intact. All the fuel tank caps were found installed on their respective fuel tank. The fuel tank caps were removed to view the inside of the fuel tanks. Fuel was not observed in any of the fuel tanks. All the fuel tank bladders were found intact, and they were not collapsed. The left side fuel selector was found on the left main fuel tank. The right side fuel selector was found on the right auxiliary fuel tank. No signs of fuel leakage were observed on the airframe or on the two engines. During the recovery operation, no usable fuel was recovered from the airplane.

According to the Cessna 340A Pilot's Operating Handbook and Federal Aviation Administration (FAA) Approved Airplane Flight Manual, the airplane has an optional fuel low level warning light system. This system provides a visual warning in the cockpit when the left main fuel tank or the right main fuel tank, or both, contain about 60 lbs (about 9.68 gallons) of fuel. The airplane was not equipped with this system, nor was it required to be.

The configuration of the airplane's fuel system did not allow the transfer of fuel from an auxiliary fuel tank to either of the main fuel tanks; auxiliary fuel could only be provided to the engine on the same side as the auxiliary fuel tank, via the fuel selector. The fuel in a locker fuel tank was fed to the associated main fuel tank when the transfer pump was activated by the pilot.

The airplane was not equipped with a digital engine monitor or a digital fuel flow indicator system.

A review of the airplane maintenance records showed that the airplane's most recent annual inspection was completed on October 3, 2023, at an airframe total time of 2,980.4 hours.

The FAA Pilot's Handbook of Aeronautical Knowledge FAA-H-8083-25C discusses fuel tank quantity gauges and states in part:

The fuel quantity gauges indicate the amount of fuel measured by a sensing unit in each fuel tank and is displayed in gallons or pounds. Aircraft certification rules require accuracy in fuel gauges only when they read "empty." Any reading other than "empty" should be verified. Do not depend solely on the accuracy of the fuel quantity gauges. Always visually check the fuel level in each tank during the preflight inspection, and then compare it with the corresponding fuel quantity indication.

The document also discusses fuel consumption rates and states in part:

The rate of fuel consumption depends on many factors: condition of the engine, propeller/rotor pitch, propeller/rotor revolutions per minute (rpm), richness of the mixture, and the percentage of horsepower used for flight at cruising speed. The pilot should know the approximate consumption rate from cruise performance charts or from experience. In addition to the amount of fuel required for the flight, there should be sufficient fuel for reserve. When estimating consumption, you must plan for cruise flight as well as startup and taxi, and higher fuel burn during climb. Remember that ground speed during climb is less than during cruise flight at the same airspeed. Additional fuel for adequate reserve should also be added as a safety measure. -'

In [210...

```
# Check for nulls
df2.ProbableCause.isnull().sum()
```

Out[210...

1057

In [211...

```
df2[df2['ProbableCause'].isnull()]
```

Out[211...

	ProbableCause	City	EventDate	Year	Month	State	FactualNarrative	FatalInjuryCount
7072	None	Lasne	2020-09-13 16:32:00+00:00	2020	9	None	The government of Belgium has notified the NTS...	2
7145	None	Teresina	2020-08-28 16:15:00+00:00	2020	8	None	The government of Brazil has notified the NTSB...	1
7158	None	Springs	2020-08-26 20:20:00+00:00	2020	8	None	The government of South Africa has notified th...	1
7250	None	Lobios (Ourense)	2020-08-08 10:20:00+00:00	2020	8	None	The Comision de Investigacion de Accidentes y ...	1
7291	None	Tarapaca-Amazonas	2020-07-30 18:03:00+00:00	2020	7	None	The government of Colombia has notified the NT...	1
...
28780	None	Fairmont	2007-10-26 19:00:00+00:00	2007	10	None	On October 26, 2007, about 1900, mountain dayl...	3
28799	None	Araya	2007-10-23 14:20:00+00:00	2007	10	None	On October 10, 2007, about 1420 Atlantic stand...	2
28816	None	Vancouver	2007-10-19 16:03:00+00:00	2007	10	None	On October 19, 2007, about 1603, pacific dayli...	1
28840	None	Bamfield	2007-10-13 15:00:00+00:00	2007	10	None	On October 13, 2007, about 1500 pacific daylig...	3
28842	None	Great Harbor	2007-10-13 10:36:00+00:00	2007	10	None	On October 13, 2007, at 1036 eastern daylight ...	1

1057 rows × 14 columns

```
In [223... df2.iloc[3]['ProbableCause']
```

```
Out[223... 'The pilot's improper fuel planning, which resulted in fuel exhaustion and a total loss  
of power to both engines.'
```

```
In [224... # Have a Look at example of missing summary reports  
df2['FactualNarrative'].loc[18111]
```

```
Out[224... 'During the landing flare to a private field, the airplane encountered a crosswind from  
the southwest. The wind pushed the airplane to the north and the airplane landed prematu  
rely on a high berm located to the north of the runway. The landing gear collapsed and t  
he airplane was damaged by the postimpact fire. The pilot did not report any mechanical  
anomalies with the airplane. -'
```

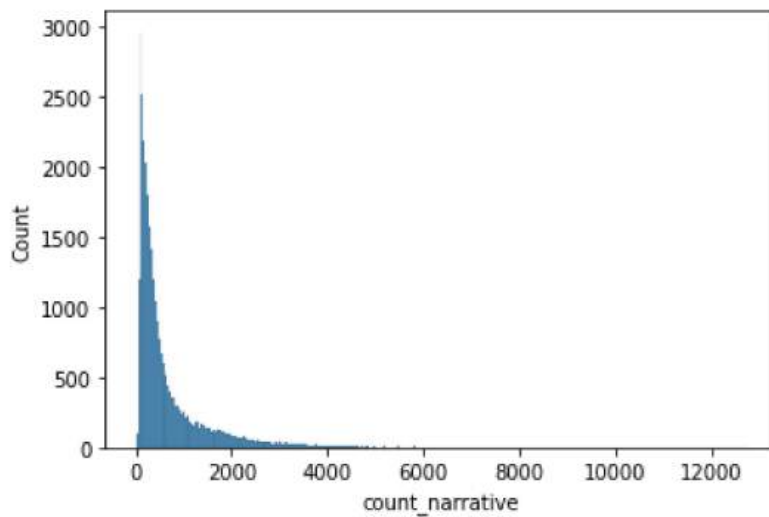
```
In [225... # Drop the missing rows  
df2 = df2[~df2['ProbableCause'].isnull()]
```

```
In [226... # Get count of the words in the report and summary text  
df2["count_narrative"] = df2['FactualNarrative'].str.split().apply(len)  
df2["count_prob_cause"] = df2['ProbableCause'].str.split().apply(len)
```

```
In [227... # Look at the distribution of the word count for the narrative  
df2['count_narrative'].describe()
```

```
Out[227... count      31003.000000  
mean         685.914428  
std          919.107312  
min           3.000000  
25%          175.000000  
50%          344.000000  
75%          784.000000  
max         12730.000000  
Name: count_narrative, dtype: float64
```

```
In [228... # Plot distribution  
sns.histplot(data = df2, x = 'count_narrative');
```



Right skewed with median length of 328 words. There are some long reports in the data of up to 12730 words which would be difficult for sentence transformer embeddings that work on text up to 256 tokens.

In [229...


```
# Look at the distribution of the word count for the summary
df2['count_prob_cause'].describe()
```

Out[229...

```
count    31003.000000
mean      25.335548
std       13.891441
min        1.000000
25%       16.000000
50%       23.000000
75%       31.000000
max       199.000000
Name: count_prob_cause, dtype: float64
```

In [219...

```
# Plot distribution
sns.histplot(data = df2, x = 'count_prob_cause');
```



Some right skew with median of 23 words. The maximum length is 199 and a lot of the important reasons for the accident are summarised in this text

Observations

- The full reports do contain a lot of information which could be useful but some of them are very long and would need to be removed for BERTopic to work efficiently.
- These reports would probably work better with LDA which handles long texts.
- For this exercise the summary of the report in the column 'Probable Cause' is taken for training the model.
- This contains much shorter succinct text which would be more helpful in running BERTopic. It will also not require the level of text cleaning needed for the longer report which contains a lot of repetitive information about the date and location, both of which are not required for this analysis as we have that information in the table data and also a lot of explanatory words which we might include as stop words and again are not information rich.
- The downside of using just BERTopic is that this throws out just one topic for a piece of text, rather than a range of probable topics which we get in LDA.
- Notebook 2 continues the work on this

In [230...

```
df2.to_csv(r'C:\Users\imoge\AllMLProjects\Data\USAAccident2001to2021Cleaned.csv')
```

In [221...

```
# Use a function to plot a wordcloud (we don't need stopwords)
def wordcloud(df,title):

    # Create text object
    text = " ".join(review for review in df3['ProbableCause'])

    # Set stopwords
    stopwords = set(STOPWORDS)
    stopwords = ["resulted"] + list(STOPWORDS)

    # Create and generate a word cloud image:
    wordcloud = WordCloud(max_words = 500,
                           colormap = 'Blues',
                           collocations = True,
                           background_color="Black",
                           stopwords = stopwords,
                           width=1000, height=600).generate(text)

    # Plot the wordcloud
    plt.figure(figsize = (10,8))
    plt.imshow(wordcloud, interpolation="bilinear")
    plt.axis("off")
    plt.title(title)
    plt.show()
```

In [222...

```
wordcloud(df2,'Accident Report Summary')
```



```
#import contractions
#from nltk.tokenize import word_tokenize
#import string
```

```
# test code for text cleaning
#a = df["FactualNarrative"].apply(lambda x: x.strip().lower())
#b = a.apply(lambda x: contractions.fix(x))
#c = b.apply(lambda x: word_tokenize(x))
#punc = string.punctuation
#stop = ['january', 'february', 'march', 'april', 'may', 'june',
        # 'july', 'august', 'september', 'october', 'november', 'december']
#d = c.apply(lambda x: [w for w in x if w not in stop])
#e = d.apply(lambda x: [word for word in x if word not in punc])
#f = e.apply(lambda x: [n for n in x if not n.isnumeric()])
#g = f.apply(lambda x: [e for e in x if e.encode("ascii", "ignore")])
#h = g.apply(lambda x: " ".join(x)) #(using stemmed output as results in reduced duplicates)
#h
```

```
#i = " ".join(w for w in h)
#i
```

```
def clean_text(df):

    df["FactualNarrative"] = df["FactualNarrative"].apply(lambda x: x.strip())

    # Expand contractions
    df["FactualNarrative"] = df["FactualNarrative"].apply(lambda x: contractions.fix(x))

    # Tokenize
    df["tok"] = df["FactualNarrative"].apply(lambda x: word_tokenize(x))

    # Remove punctuation
```

```

punc = string.punctuation
df["punct"] = df["tok"].apply(lambda x: [word for word in x if word not in punc])

# Stopwords removal
stop = ['January', 'February', 'March', 'April', 'May', 'June', 'July', 'August', 'Septembe
df["stop"] = df["punct"].apply(lambda x: [w for w in x if w not in stop])

# Remove numbers, except words that contain numbers.
df["num"] = df["stop"].apply(lambda x: [n for n in x if not n.isnumeric()])

# Remove non ascii characters
df["ascii"] = df["num"].apply(lambda x: [e for e in x if e.encode("ascii", "ignore")

# Join the text into strings and join the strings into one text
df["joined"] = df["ascii"].apply(lambda x: " ".join(x)) #(using stemmed output as re
text = " ".join(e for e in df.joined)

return text

```

In [117...

```

# Get count of the words in the report text
df3["count"] = df3['FactualNarrative'].str.split().apply(len)
df3['count'].describe()

```

Out[117...

```

count      31003.000000
mean         685.914428
std          919.107312
min           3.000000
25%          175.000000
50%          344.000000
75%          784.000000
max         12730.000000
Name: count, dtype: float64

```

In [118...

```

# How many words overall
df3['count'].sum()

```

Out[118...

```

21265405

```